HPV Whole-Genome Nanopore Analysis Pipeline

Pipeline Report & User Manual — shared cluster deployment

Bright Adu, PhD

2026-06-06

Contents

Document control **3**

1 Scope **4**

2 Executive summary **5**

3 System architecture **6**
3.1 Shared-pool layout (all-users model) 6
3.2 Stage-flow diagram 7

4 The six conda environments **8**
4.1 hpv-qc — read QC, filtering, coverage (Python 3.13.13) 8
4.2 hpv-artic — ARTIC consensus + Clair3, CPU (Python 3.11.15) 9
4.3 hpv-asm — optional de novo path (Python 3.12.13) 10
4.4 hpv-py — typing/binning, summaries, figures (Python 3.13.13) 10
4.5 hpv-typing — optional k-mer / competitive read typing (Python 3.14.5) 11
4.6 hpv-smk — Snakemake workflow manager (Python 3.13.13) 11

5 The analysis pipeline in detail **13**
5.1 Stage 1 — Ingest 13
5.2 Stage 2 — Read QC 13
5.3 Stage 3 — Amplicon-aware length filter 13
5.4 Stage 4 — Type assignment (binning) 13
5.5 Stage 5 — Per-type consensus (ARTIC + Clair3, with junction handling) 14
5.6 Stages 6–7 — Coverage QC, summary & figures 14

6 Origin-junction handling **15**
6.1 Why there is a seam 15
6.2 Why a single rotated frame is not enough 15
6.3 The implemented two-run solution 15
6.4 Evidence it recovers the seam (`junction_proof.tsv`) 16

7 Clair3 model selection **17**
7.1 Installed models 17
7.2 Auto-selection and override 17

8 Running a real analysis **18**

8.1 Running as any user (zero setup) . . . 19
8.2 Distribution model (source vs deployed) . . . 19
8.3 Before you start — copy these off the instrument . . . 20
8.4 The sample sheet . . . 20
8.5 Equivalent: Snakemake . . . 20
8.6 Equivalent: Nextflow . . . 20
8.7 SLURM / scheduler batch jobs . . . 21
8.8 Where outputs land . . . 21

9 Interpreting outputs **22**
9.1 run_summary.tsv . . . 22
9.2 consensus_qc.tsv . . . 22
9.3 Consensus FASTAs / VCFs . . . 22
9.4 The QC figures (`results/figures/`) . . . 23
9.5 Design caveats analysts must expect . . . 23

10 Validation record **26**
10.1 (1) End-to-end on the real ARTIC + Clair3 path . . . 26
10.2 (2) All-users deployment verified under a non-admin account . . . 27

11 Troubleshooting & FAQ **28**

12 Appendix **30**
12.1 Core runtime version strings . . . 30
12.2 Full `conda list` per env . . . 30
12.3 The all-users permissions model . . . 31
12.4 Design provenance . . . 31
12.5 How to report issues . . . 31

Document control

Field	Value
Document	HPV Whole-Genome Nanopore Analysis Pipeline — Report & User Manual
Author	Bright Adu, PhD — Genomics and Bioinformatics Core, Noguchi Memorial Institute for Medical Research, University of Ghana, Legon, Ghana · ORCID 0000-0002-3622-2137
Manual version	**v1.1** (final; deployed-shared-install reality)
Date	**2026-06-06**
Scheme version	Primer scheme **v5** (canonical accessions; 12 high-risk HPV types)
Shared install root	`INSTALL_ROOT=/hpc/opt/hpv-bioinfo`
Shared conda home	`HPV_CONDA_HOME=/hpc/opt/conda` (named `hpv-*` envs)
Deployed pipeline	`HPV_PIPELINE=/hpc/opt/hpv-bioinfo/pipeline` (run via `hpv-run`)
Validation (callers)	`artic+clair3`, Clair3 model `r1041_e82_400bps_hac_g615`
Validation (deployment)	`verify_env.sh` under a non-admin account: **53 passed, 0 failed**

How this document was produced — read this first. Every fact here was pulled from the live system or actual run outputs, never from memory:

- **Tool versions** are the actual installed builds, read from a live `conda list` of all six shared envs (`/hpc/opt/conda/envs/hpv-{qc,artic,asm,py,typing,smk}`), captured on the production cluster and shipped beside this manual as **`hpv_env_versions.txt`**. Runtime `--version` banners (e.g. `minimap2 2.31-r1302`, `nextflow 26.04.3`) are quoted where they differ from the conda package string.
- **Validation evidence** (engine, model, %genome, seam N, co-infection, and the non-admin `verify_env.sh` result) is quoted from the actual run outputs.

Chapter 1

Scope

This document is both a **reference report** (what the pipeline is, how it is built, and how it was validated) and a **practical user manual** (how any lab member runs it on real sequencing data). It covers the HPV **whole-genome tiling-amplicon** Nanopore analysis system for **12 high-risk HPV types — HPV 16, 18, 31, 33, 35, 39, 45, 52, 56, 58, 59, 66** — using a lineage-robust primer scheme (**v5**) designed for sub-Saharan African genomic diversity. Each type is amplified as a **6-amplicon, dual-pool** tiling set spanning the full ~7.8–7.9 kb circular genome, including the **origin-spanning junction amplicon**. The platform is **CPU-only** and runs on the shared HPC cluster; basecalling and demultiplexing happen upstream on the instruments (GridION / PromethION P2 Solo). The pipeline takes demultiplexed `fastq_pass/barcodeNN/` reads and produces per-type whole-genome consensus sequences, variant calls, coverage QC, and publication-quality figures.

Audience: all lab users on the cluster who run analyses, plus a methods-section reader who needs the provenance and validation record.

Chapter 2

Executive summary

The system performs **end-to-end HPV whole-genome reconstruction from Nanopore tiling-amplicon reads**:

1. **Ingest** demultiplexed `fastq_pass/barcodeNN/` reads per sample.
2. **Read QC** (length/quality stats; optional pycoQC run report from `sequencing_summary`).
3. **Amplicon-aware length filter** (retains the short ~1.1 kb HPV52 L1 gap-fill amplicon while removing chimeras / short fragments).
4. **Type assignment (binning)** by competitive minimap2 mapping against all 12 references — **co-infection-aware** (a sample can resolve to more than one type).
5. **Per-type consensus** via the real **ARTIC `minion` + Clair3** path, including the **two-run origin-junction handling** that recovers the seam.
6. **Coverage QC** (per-amplicon depth, pool balance, completeness).
7. **Summary tables + publication figures**.

Key properties:

- **CPU-only.** No GPU and no basecaller on the cluster — Clair3 and everything else run on CPU. Basecalling happens on the **GridION** and **PromethION P2 Solo** (both R10.4.1, 400 bps chemistry).
- **Shared, all-users deployment.** Named `hpv-*` conda envs live in the shared cluster conda (`/hpc/opt/conda/envs`); reference data, models, schemes **and the pipeline code itself** live on the shared pool (`/hpc/opt/hpv-bioinfo`; code at `$HPV_PIPELINE`, run via `hpv-run`). The login profile auto-loads variables and PATH silently, so **a non-admin user needs no setup, no GitHub clone, and nothing in their home directory to run the pipeline.** All run-time output goes to the user's own working directory; nothing is written to the shared install.
- **Validated two ways (see the Validation record).** (1) End-to-end on the real **ARTIC 1.10.1 + Clair3 2.0.1** path (engine `artic+clair3`, 100 % genome, 0 % N, seam N = 0, co-infection split). (2) The all-users deployment verified by `verify_env.sh` under a genuine **non-admin** account: **53 passed, 0 failed**.

  Software validation vs biological validation. Both proofs are *software* validations — they show the pipeline runs end to end and deploys correctly for all users. **Biological validation on control material (characterised isolates / clinical positives) is the next, separate, wet-lab step.**

Chapter 3

System architecture

3.1 Shared-pool layout (all-users model)

Two shared locations, both on `/hpc` so every node sees them identically:

Component	Location	Notes
Named conda envs	`/hpc/opt/conda/envs/hpv-{qc,artic,asm,py,typing,smk}`	Created like the lab's `wgs_*` envs; resolve cluster-wide by name
Shared data/tools	`INSTALL_ROOT=/hpc/opt/hpv-bioinfo`	Owned `root:bioinfo`, dirs 2775 setgid, files world-readable
Primer schemes	`$SCHEME_DIR` = `…/primer-schemes/HPV*/V5/`	`reference.fasta`, `primer.bed`, insert/amplicon BEDs
References	`$REF_DIR` = `…/references/hpv_all_types.fasta`	+ `.fai`, + `.mmi` minimap2 index
Clair3 models	`$MODELS_DIR` = `…/clair3_models/`	r10.4.1 400 bps **fast / hac / sup** (read-only)
Binaries	`…/bin`	On PATH; holds `nextflow` and the `hpv-run` launcher
Pipeline code	`$HPV_PIPELINE` = `…/pipeline`	Deployed `run_pipeline.py`, `hpvpipe.py`, `qc_report.py`, `qc_lib.py`, `simulate_test_data.py`, `Snakefile`, `main.nf`, `nextflow.config`, `testdata/` — the copy **all users run**

Component	Location	Notes
Manual / docs	…/docs	This manual (.md/.pdf/.docx) + figures
Profile	/etc/profile.d/hpv-bioinfo.sh	Auto-loads SCHEME_DIR, REF_DIR, MODELS_DIR, HPV_PIPELINE, HPV_ARTIC_RUN; provides hpv-activate, hpv-run
Logs / scheme summary	…/logs/ ; scheme_summary.tsv	Install logs + per-type scheme provenance

The all-users rule (HARD): the shared install is read/execute for everyone, write for admins only. No analysis step writes to /hpc/opt/hpv-bioinfo or the shared conda. All outputs, caches, work and temp go to the user's CWD/HOME — the profile points NXF_HOME, APPTAINER_CACHEDIR and CONDA_PKGS_DIRS into $HOME, and Clair3 models are read-only (passed by path, never written back).

3.2 Stage-flow diagram

```
Instrument (GridION / PromethION P2 Solo): basecalling + demultiplexing
   |
   v   fastq_pass/barcodeNN/*.fastq.gz  (+ sequencing_summary_*.txt)
   |
 [1] INGEST         concat per-barcode FASTQ -> per-sample        (hpv-py)
 [2] READ QC        length/quality stats (+pycoQC if seq-summary) (hpv-qc)
 [3] LENGTH FILTER  amplicon-aware window (keeps HPV52 ~1.1 kb)   (hpv-qc)
 [4] TYPE ASSIGN    competitive minimap2 vs 12 refs; co-infection (hpv-py)
 [5] CONSENSUS      ARTIC minion + Clair3                        (hpv-artic)
       main:     5 non-wrapping amplicons on the canonical ref
       junction: amp3 on a ROTATED ref -> rotate back + merge seam
 [6] COVERAGE QC    per-amplicon depth, pool balance             (hpv-py)
 [7] SUMMARY+FIGS   run_summary / consensus_qc / junction_proof  (hpv-py)
   |
   v
 results/ : consensus FASTA, VCF, coverage, figures, summary -> USER WORKDIR
```

The same graph is implemented three ways with identical stages: run_pipeline.py (single-driver, recommended), Snakefile (Snakemake), and main.nf (Nextflow). The core stage logic lives in pipeline/hpvpipe.py.

Chapter 4

The six conda environments

All six are **named envs in the shared cluster conda** and activate cluster-wide:

```
conda activate hpv-qc        # exactly like: conda activate wgs_prokka
# or, equivalently:
hpv-activate qc
```

> **Version tables carry the actual installed builds.** Each shows the spec **floor** (>= pin from `server-setup/envs/<env>.yml`) and the **installed (live)** build read from `conda list` on the cluster (full capture in `hpv_env_versions.txt`). The Python interpreter version is in each heading. Re-confirm on the node with `conda list -n hpv-<name>` or `conda run -n hpv-<name> <tool> --version`.

4.1 hpv-qc — read QC, filtering, coverage (Python 3.13.13)

Purpose: read-level QC, length/quality filtering, coverage, generic utilities. **Use it for:** read QC and the length filter, plus ad-hoc `seqkit`/`samtools`/`mosdepth` inspection.

Tool	Floor	Installed (live)
python	3.10	3.13.13
nanoplot	1.42	1.46.2
nanocomp	1.23	1.25.6
pycoqc	2.5	2.5.0.3
multiqc	1.21	1.35
fastqc	0.12	0.12.1
seqkit	2.6	2.13.0
csvtk	0.30	0.37.0
mosdepth	0.3.8	0.3.14
bedtools	2.31	2.31.1
samtools	1.19	1.23.1
chopper	0.7	0.13.0
htslib · tabix · pigz	—	1.23.1 · 1.11 · 2.8

```
hpv-activate qc
```

```
NanoPlot --fastq fastq_pass/barcode01/*.fastq.gz -o qc/barcode01 --N50
pycoQC -f sequencing_summary_*.txt -o qc/pycoQC.html   # needs the seq
↪ summary
seqkit stats -a fastq_pass/barcode01/*.fastq.gz
```

4.2 hpv-artic — ARTIC consensus + Clair3, CPU (Python 3.11.15)

Purpose: the core variant-calling / consensus environment — **ARTIC fieldbioinformatics** with **Clair3 2.x** (PyTorch models), CPU-only. **Use it for:** per-type consensus. The pipeline shells into this env via `HPV_ARTIC_RUN="conda run -n hpv-artic"`.

Tool	Floor	Installed (live)
python	3.9	3.11.15
artic (fieldbioinformatics)	1.6	**1.10.1** (engine of record)
clair3	2.0.1	**2.0.1**
pytorch-cpu (Clair3 backend)	—	**2.10.0** (`cpu_generic` — confirms CPU-only)
align_trim	1.1.0	1.1.0
primalbedtools	1.0.0	1.0.0
minimap2	2.26	2.31 (runtime `2.31-r1302`)
samtools	1.19	1.23.1
bcftools	1.19	1.23.1
htslib	1.19	1.23.1
seqtk · biopython · pyfaidx · tabix	—	1.5 · 1.87 · 0.9.0.4 · 1.11

Also present: `whatshap` 2.8, `longphase` 2.0.2 (phasing), bwa 0.7.19, `mafft` 7.526, `pysam` 0.24.0.

> The ARTIC `minion.log.txt` from the validation run confirms the live tool chain: `minimap2 -a -x map-ont`; `align_trim --normalise 100 … --primer-match-threshold 35 --min-mapq 20`; and `run_clair3.sh … --platform=ont --model_path=…/r1041_e82_400bps_hac_g615`.

```
hpv-activate artic
artic --version
# ARTIC offline scheme loading uses the local files directly (no manifest
↪ fetch):
artic minion --threads 16 \
  --bed "$SCHEME_DIR/HPV16/V5/HPV16.primer.bed" \
  --ref "$SCHEME_DIR/HPV16/V5/HPV16.reference.fasta" \
  --model-dir "$MODELS_DIR" \
  SAMPLE_X.HPV16     # the pipeline orchestrates this — see Origin-junction
↪ handling
```

4.3 hpv-asm — optional de novo path (Python 3.12.13)

Purpose: optional de novo assembly / polishing for divergent, novel, or hard co-infection cases. **`medaka` here is the standalone polisher only — NOT the ARTIC variant caller** (variant calling is Clair3 in `hpv-artic`).

Tool	Floor	Installed (live)
python	3.9	3.12.13
flye	2.9.3	2.9.6
medaka	1.11	**2.2.2**
racon	1.5	1.5.0
minimap2	2.26	2.31 (runtime `2.31-r1302`)
samtools	1.19	1.23.1
bcftools	—	1.23.1
any2fasta	—	0.8.1
pytorch-cpu (medaka backend)	—	2.9.1 (`cpu_mkl`)

```
hpv-activate asm
flye --nano-hq typing/SAMPLE_X.HPV16.fastq.gz \
  --out-dir asm/SAMPLE_X.HPV16 --threads 16
medaka_consensus -i typing/SAMPLE_X.HPV16.fastq.gz \
  -d draft.fasta -o medaka_out -t 16
```

4.4 hpv-py — typing/binning, summaries, figures (Python 3.13.13)

Purpose: the Python analysis environment — type assignment (binning), summary tables, and **publication QC figures**. **Use it for:** ingest, binning, coverage QC, figures, and to drive the whole pipeline (`run_pipeline.py` / `qc_report.py`).

Tool	Floor	Installed (live)
python	3.10	3.13.13
pysam	0.22	0.24.0
biopython	1.83	1.87
pandas	2.1	3.0.3
numpy	1.26	2.4.6
scipy	1.11	1.17.1
pyfaidx	0.7	0.9.0.4
scikit-bio	0.5.9	0.7.3
matplotlib	3.8	3.10.9
seaborn	0.13	0.13.2
plotly	5.20	6.8.0
python-kaleido	0.2.1	1.3.0
minimap2	2.26	2.31 (runtime `2.31-r1302`)
samtools	1.19	1.23.1
mosdepth	0.3.8	0.3.14

Tool	Floor	Installed (live)
bedtools	2.31	2.31.1
fonttools	—	4.63.0

Also present: `statsmodels` 0.14.6.

```
hpv-activate py
hpv-run --help    # == python "$HPV_PIPELINE/run_pipeline.py" --help
                  # hpv-run runs the orchestrator in THIS env (hpv-py) and
                   ↪  calls other
                  # envs per-stage via conda run — see "Running a real
                   ↪  analysis".
# regenerate figures from an existing results/ dir:
PYTHONPATH="$HPV_PIPELINE" \
  python "$HPV_PIPELINE/qc_report.py" \
  --results results --scheme-dir "$SCHEME_DIR"
# prove the figure module on synthetic demo data:
PYTHONPATH="$HPV_PIPELINE" \
  python "$HPV_PIPELINE/qc_report.py" --demo --outdir /tmp/demo/figures
```

4.5 hpv-typing — optional k-mer / competitive read typing (Python 3.14.5)

Purpose: optional k-mer / competitive-mapping read typing against a custom 12-type HPV database, plus BLAST / Entrez utilities for ad-hoc lookup. **Use it for:** an independent typing cross-check (the primary binning runs in `hpv-py`).

Tool	Floor	Installed (live)
python	3.10	3.14.5
minimap2	2.26	2.31 (runtime `2.31-r1302`)
samtools	1.19	1.23.1
pysam	0.22	0.24.0
kraken2	2.1.3	2.17.1
seqkit	2.6	2.13.0
blast	—	2.17.0
entrez-direct	—	24.0

```
hpv-activate typing
minimap2 -x map-ont -a "$REF_DIR/hpv_all_types.fasta" reads.fastq.gz \
  | samtools sort -o aln.bam
# kraken2 needs a custom 12-type HPV DB built separately by an admin
```

4.6 hpv-smk — Snakemake workflow manager (Python 3.13.13)

Purpose: the Snakemake workflow engine (alternative to the Nextflow binary on PATH).

Tool	Floor	Installed (live)
python	3.10	3.13.13
snakemake-minimal	8.10	**9.22.0**
pandas	2.1	3.0.3

The installed Snakemake is the **9.x** series (floor was 8.10). The example commands use options stable across 8/9.

Chapter 5

The analysis pipeline in detail

Each stage lists what it does, the env/tool, inputs/outputs, and key parameters. Defaults are read from `run_pipeline.py` / `hpvpipe.py`.

5.1 Stage 1 — Ingest

- **Env/tool:** `hpv-py` (pure Python / gzip).
- **In:** `fastq_pass/<barcode>/*.fastq{,.gz}` per the sample sheet. **Out:** `results/reads/<sample>.fastq.gz` (one concatenated FASTQ per sample).
- Maps each `barcodeNN` to its `sample_id` from `sample_sheet.csv`.

5.2 Stage 2 — Read QC

- **Env/tool:** `hpv-qc` (NanoPlot/pycoQC); the driver also writes per-read length tables.
- **In:** per-sample reads (+ optional `sequencing_summary_*.txt`). **Out:** `results/qc/read_stats.tsv`, `results/qc/lengths/<sample>.lengths.tsv`, and a **pycoQC** run report *only if* the `sequencing_summary` was uploaded.

5.3 Stage 3 — Amplicon-aware length filter

- **Env/tool:** `hpv-qc` (`chopper` if present; else a built-in AWK/Python filter).
- **In:** per-sample reads. **Out:** `results/filtered/<sample>.filt.fastq.gz`.
- **Key parameters:** `--min-len 900` (default), `--max-len 4800` (default), quality `-q 9`.

  Caveat — the length window must keep HPV52 amp7. The HPV52 L1 gap-fill amplicon (**amp7, ~1117 bp**) is *shorter* than the 2–4 kb tiling amplicons. The minimum length must stay **≥ ~900 bp** or this amplicon is filtered out and HPV52 loses L1 coverage.

5.4 Stage 4 — Type assignment (binning)

- **Env/tool:** `hpv-py` — competitive `minimap2` mapping (`hpvpipe.py bin`).

- **In:** filtered reads + `references/hpv_all_types.fasta`. **Out:** `results/typing/<sample>.types.tsv` per-type `results/typing/<sample>.<TYPE>.fastq.gz`, and `results/typing/type_counts.tsv`.
- **Key parameters:** `--min-reads 20` (default; a type must reach this read count to be "present"), `--min-ident 0.80` (default minimum identity for read assignment), `--threads`.
- **Co-infection handling:** each read is assigned to its best-scoring type; **a sample can yield multiple present types** (e.g. the validated HPV16/HPV58 co-infection). Every present type is carried independently through consensus.

5.5 Stage 5 — Per-type consensus (ARTIC + Clair3, with junction handling)

- **Env/tool:** `hpv-artic` (ARTIC `minion -> align_trim -> Clair3`), orchestrated by `hpvpipe.py consensus`; the portable pysam pileup engine is an explicit opt-in fallback.
- **In:** per-type reads + per-type scheme (`primer.bed`, `reference.fasta`) + Clair3 models. **Out:** `results/consensus/<sample>.<TYPE>.consensus.fasta`, VCFs, per-amplicon depth tables, `…consensus_qc.tsv`, `…junction_proof.tsv`.
- **Key parameters:** `--use-artic` (real ARTIC+Clair3 path), `--require-artic` (strict: abort on any failure), `--allow-fallback` (opt in to the portable engine), `--mindepth 3` (default; positions below become N), `--model`/`--model-dir` (Clair3 model — see Clair3 model selection). ARTIC internally runs `align_trim --normalise 100 --primer-match-threshold 35 --min-mapq 20`.
- The **dual-pool** design (Pool 1 / Pool 2 alternating amplicons) is honoured by ARTIC's read-group handling; the origin-junction two-run mechanism is detailed in the next section.

5.6 Stages 6–7 — Coverage QC, summary & figures

- **Env/tool:** `hpv-py` (`qc_report.py`, mosdepth-style per-base/amplicon coverage).
- **Out:** `results/coverage/*`, `results/run_summary.tsv`, `results/consensus/consensus_qc.tsv`, and `results/figures/` (every figure as **SVG + PDF + 300-dpi PNG**, colour-blind-safe, embedded fonts). See Interpreting outputs.

Chapter 6

Origin-junction handling

This is the system's most distinctive feature.

6.1 Why there is a seam

HPV genomes are **circular**, but reference FASTAs and ARTIC are **linear**. To tile the whole genome the scheme must place **one amplicon that spans the origin** (position 0) — this is **amp3** in every type (`scheme_summary.tsv`: `junction_amplicon = 3`, `junction_wraps_origin = YES` for all 12 types). On a linear reference, that amplicon's right primer coordinate is *smaller* than its left (it wraps past the end back to the start). For HPV16, `rotation.json` records amp3 as canonical **LEFT = [6722, 6744], RIGHT = [874, 896]** on a genome of length **L = 7906** — i.e. it runs from ~6722 across the origin to ~896.

ARTIC's `align_trim` computes each amplicon's length as `right_end - left_start`. For a wrapping amplicon that is **negative**, and `align_trim` then attempts `np.zeros(length)` — historically this crashed with *"ValueError: negative dimensions are not allowed."* The seam must therefore be handled explicitly, or amp3 (and the region across the origin) is lost.

6.2 Why a single rotated frame is not enough

The intuitive fix — rotate the reference so the origin moves to the middle and amp3 becomes contiguous — works **for amp3**, but it just **moves the wrap to whichever amplicon now spans the new origin**. A single rotated frame always leaves *some* amplicon wrapping. There is no single linear frame in which all six amplicons of a circular tiling scheme are simultaneously non-wrapping.

6.3 The implemented two-run solution

The pipeline runs ARTIC **twice, each on a frame where every amplicon it sees is positive-length**, then merges:

1. **Main run — canonical reference, amp3 removed.** ARTIC is given a primer BED with the junction amplicon excluded (`<sample>.<TYPE>.no_amp3.primer.bed`), so `align_trim`

only ever sees the **5 non-wrapping amplicons**. These are mapped, primer-trimmed, and called by Clair3 normally. The origin seam is left as `N` in this consensus.

2. **Junction run — rotated reference, amp3 only.** A per-type **rotated reference** is built (`rotate–build`): the origin is shifted by an offset (~L/2) so amp3 becomes one **contiguous, positive-length** amplicon. For HPV16, `rotation.json` shows amp3 rotated to **LEFT = [2769, 2791], RIGHT = [4827, 4849]** (offset **3953**) — clearly positive. ARTIC + Clair3 call amp3 on this frame.
3. **Rotate back & merge.** The amp3 consensus is mapped back to canonical coordinates (`pos_canon = (pos_rot + offset) mod L`) and used to **fill only the seam positions still `N`** in the main consensus. Length is unchanged; no double-counting.

Both runs go through the real `align_trim` + Clair3, and **neither run ever passes a negative-length amplicon** to `align_trim`. The validation-run logs show exactly this: the main run calls `align_trim … <sample>.no_amp3.primer.bed`, and the junction run (`*.artic.rotj`) calls `align_trim … HPV16.rot.primer.bed` on `HPV16.rot.reference.fasta`.

6.4 Evidence it recovers the seam (`junction_proof.tsv`)

The per-(sample×type) `junction_proof.tsv` records the merge. From the validation run:

metric	HPV16 (SAMPLE_COINF)	HPV58 (SAMPLE_COINF)
junction_amplicon	amp3	amp3
junction_engine	artic+clair3	artic+clair3
rotation_offset	3953	3912
seam_region_bp	2036	2208
canonical_only_seam_N	448	588
merged_seam_N	**0**	**0**
seam_bases_filled_from_rotated	448	588
seam_identity_to_reference_pct	100.0	100.0
continuous_across_origin	YES	YES
length_unchanged	YES	YES

The canonical-only consensus would have left **448 (HPV16) / 588 (HPV58)** `N` bases across the origin; after the rotated amp3 run and merge, **merged seam N = 0** at **100 % identity to reference**, continuous across the origin, length unchanged.

Chapter 7

Clair3 model selection

7.1 Installed models

`$MODELS_DIR` (= `/hpc/opt/hpv-bioinfo/clair3_models/`) holds the full R10.4.1 **400 bps** Clair3 2.x set — **fast / hac / sup** — because the lab runs both a GridION and a PromethION P2 Solo and the Dorado preset can differ per run. The `hac` build used in validation is `r1041_e82_400bps_hac_g615`; all three tiers are installed (the setup step fails loudly if any of fast/hac/sup is missing). List them on the node with `ls "$MODELS_DIR"`.

> **Clair3 runs on CPU.** In `hpv-artic`, Clair3 2.0.1 is backed by **`pytorch-cpu 2.10.0`** (`cpu_generic`) — no CUDA/GPU build is installed. Inference uses the PyTorch (`.pt`) models under `$MODELS_DIR`.

7.2 Auto-selection and override

- **Auto:** the pipeline reads the FASTQ header `basecall_model_version_id` and selects the matching installed model by chemistry/speed/tier; it warns and uses the nearest installed model if an exact match is absent.
- **Override:** force a model explicitly with `--model <name> --model-dir "$MODELS_DIR"`.

```
# inspect what model selection would pick for a given FASTQ:
python "$HPV_PIPELINE/hpvpipe.py" model-select \
  --reads results/typing/SAMPLE_X.HPV16.fastq.gz \
  --models-dir "$MODELS_DIR"

# GridION header  …_400bps_hac@v5.2.0  -> selects r1041_e82_400bps_hac_*
# P2-Solo header  …_400bps_sup@v5.0.0  -> selects r1041_e82_400bps_sup_*
# force a tier if needed (passed through to hpvpipe consensus):
#   --model r1041_e82_400bps_sup_g615 --model-dir "$MODELS_DIR"
```

Chapter 8

Running a real analysis

> **How `hpv-run` uses the six environments — it does *not* activate them all.** `hpv-run` launches the orchestrator (`run_pipeline.py`) inside **one** environment, **`hpv-py`**, via `conda run -n hpv-py`. You never manually activate anything. The pipeline's own logic (ingest, binning, coverage QC, figures) is Python and runs in `hpv-py`. A stage that needs a *different* environment is invoked as a **subprocess** via `conda run -n <env>` — it is never activated in your shell. Specifically, the consensus stage calls **`hpv-artic`** through `HPV_ARTIC_RUN="conda run -n hpv-artic …"` to run `artic minion` + Clair3. Think of `hpv-py` as the conductor: it calls another env per-stage for its specific tools and then exits it, so **only `hpv-py` is ever active in the session.** The six envs are deliberately isolated (their tool/dependency versions would conflict if merged), so per-stage `conda run` is the correct way to compose them — not simultaneous activation.

Which environment runs what (default `run_pipeline.py` path):

Stage(s)	Environment	How it is invoked
Orchestration, ingest, binning, coverage QC, figures (and the portable consensus engine when `--use-artic` is off)	`hpv-py`	the home env (`conda run -n hpv-py`) — what `hpv-run` starts
Consensus: `artic minion` + Clair3 (with `--use-artic`)	`hpv-artic`	subprocess via `HPV_ARTIC_RUN="conda run -n hpv-artic …"`

(The optional read-QC tools in `hpv-qc` and the optional de novo path in `hpv-asm` are called the same way — `conda run -n <env>` — when you use them, but they are not part of the default automated `run_pipeline.py` run.) The **Explicit equivalent** block below (`conda activate hpv-py; python "$HPV_PIPELINE/run_pipeline.py" …`) *is* what `hpv-run` does internally — `hpv-py` is the home env.

8.1 Running as any user (zero setup)

The pipeline is **deployed on the shared pool** and the login profile **auto-loads** the environment **silently** at login. So a brand-new non-admin user, on any node, just logs in and runs — **no `source`, no GitHub clone, no `server-setup`, nothing in `$HOME`.** After login, `SCHEME_DIR`, `REF_DIR`, `MODELS_DIR`, `HPV_PIPELINE` and `HPV_ARTIC_RUN` are already set and `hpv-run` is already on `PATH`.

```
# log in, then go to your own working directory (outputs land here, never in
 ↪  /hpc/opt):
cd ~/my_run_2026-06

# simplest: one command, real ARTIC + Clair3 path (strict)
hpv-run \
  --fastq-pass  fastq_pass \
  --samplesheet sample_sheet.csv \
  --scheme-dir  "$SCHEME_DIR" \
  --ref         "$REF_DIR/hpv_all_types.fasta" \
  --models-dir  "$MODELS_DIR" \
  --outdir      results \
  --use-artic --require-artic
```

`hpv-run` runs `run_pipeline.py` inside the `hpv-py` env via `conda run` (no manual activation), and sets `HPV_ARTIC_RUN` so `--use-artic` reaches `hpv-artic` (see the call-out at the top of this chapter on how the environments are composed). Explicit equivalent — this is exactly what `hpv-run` does internally:

```
conda activate hpv-py
python "$HPV_PIPELINE/run_pipeline.py" \
  --fastq-pass  fastq_pass \
  --samplesheet sample_sheet.csv \
  --scheme-dir  "$SCHEME_DIR" \
  --ref         "$REF_DIR/hpv_all_types.fasta" \
  --models-dir  "$MODELS_DIR" \
  --outdir      results \
  --threads     16 \
  --use-artic --require-artic
```

Flag notes: **`--use-artic`** turns on the ARTIC `minion` + Clair3 path; **`--require-artic`** is strict (aborts non-zero on any failure — use it for any reportable result); replace it with **`--allow-fallback`** to permit the portable engine for quick triage; length/typing thresholds are tunable (`--min-len 900 --max-len 4800 --min-reads 20`).

Threading on the 72-core node: Clair3 is the bottleneck. Prefer moderate threads per job × several jobs in parallel over maxing one job. A good default is `--threads 16` with multiple samples concurrent via Snakemake/Nextflow; keep total cores and RAM within budget.

8.2 Distribution model (source vs deployed)

- **GitHub repo** — source of truth / version control (developers only).

- **`/hpc/opt/hpv-bioinfo/pipeline`** (`$HPV_PIPELINE`) — the deployed copy **everyone runs**, via `hpv-run` or `$HPV_PIPELINE/...`.
- **Updating** = an admin pulls new code and re-runs the deploy step of `setup_hpv_nanopore_env.sh`, which refreshes `$HPV_PIPELINE` in place. Ordinary users never touch `server-setup` or GitHub to run the pipeline, and have nothing to update on their side.

8.3 Before you start — copy these off the instrument

1. The **demultiplexed `fastq_pass/`** folder (`barcodeNN/` subfolders of `*.fastq.gz`).
2. **`sequencing_summary_*.txt`** — needed for the pycoQC run-health plots (channel activity, yield-over-time, pore occupancy). Without it those panels are skipped (read-derived QC still works).
3. Keep **both Pool 1 and Pool 2** reads — the pool-balance figure needs both.

Input layout the pipeline expects:

```
fastq_pass/
  barcode01/*.fastq.gz
  barcode02/*.fastq.gz
sequencing_summary_*.txt        # optional, enables pycoQC
sample_sheet.csv
```

8.4 The sample sheet

`sample_sheet.csv` — header `barcode,sample_id,expected_types,notes`:

```
barcode,sample_id,expected_types,notes
barcode01,SAMPLE_001,,
barcode02,SAMPLE_002,HPV16;HPV58,suspected co-infection
```

`barcode` = the MinKNOW `barcodeNN` folder; `expected_types` is optional and `;`-separated.

8.5 Equivalent: Snakemake

```
conda activate hpv-py
snakemake -s "$HPV_PIPELINE/Snakefile" -j4 \
  --config scheme_dir="$SCHEME_DIR" \
           ref="$REF_DIR/hpv_all_types.fasta" \
           models_dir="$MODELS_DIR" \
           fastq_pass=fastq_pass \
           samplesheet=sample_sheet.csv \
           use_artic=true threads=16
```

8.6 Equivalent: Nextflow

```
nextflow run "$HPV_PIPELINE/main.nf" \
  --fastq_pass fastq_pass \
```

```
  --samplesheet sample_sheet.csv \
  --scheme_dir "$SCHEME_DIR" \
  --ref "$REF_DIR/hpv_all_types.fasta" \
  --models_dir "$MODELS_DIR" \
  --outdir results --use_artic
# tune CPU/RAM per process in nextflow.config
```

8.7 SLURM / scheduler batch jobs

`/etc/profile.d/hpv-bioinfo.sh` auto-loads for interactive **login** shells, but batch / non-login jobs do **not** read `/etc/profile.d`. A job script must load the environment explicitly at the top:

```
#!/bin/bash
#SBATCH --job-name=hpv --cpus-per-task=16 --mem=32G
source /etc/profile.d/hpv-bioinfo.sh   # or: source
 ↪  /hpc/opt/hpv-bioinfo/etc/hpv-bioinfo.env
cd "$SLURM_SUBMIT_DIR"
hpv-run \
  --fastq-pass  fastq_pass \
  --samplesheet sample_sheet.csv \
  --scheme-dir  "$SCHEME_DIR" \
  --ref         "$REF_DIR/hpv_all_types.fasta" \
  --models-dir  "$MODELS_DIR" \
  --outdir      results \
  --threads     "$SLURM_CPUS_PER_TASK" \
  --use-artic --require-artic
```

8.8 Where outputs land

Everything is written under your **`--outdir` (default `./results`) in your own working directory** — never to `/hpc/opt`:

```
results/
  reads/<sample>.fastq.gz                          # ingested per-sample reads
  qc/read_stats.tsv, qc/lengths/                   # read QC
  filtered/<sample>.filt.fastq.gz                  # length-filtered reads
  typing/<sample>.types.tsv                        # type calls (+ per-type
 ↪  FASTQ)
  consensus/<sample>.<TYPE>.consensus.fasta        # per-type whole-genome
 ↪  consensus
  consensus/<sample>.<TYPE>.*.vcf(.gz)             # Clair3 variant calls
  consensus/<sample>.<TYPE>.consensus_qc.tsv       # per-type QC row
  consensus/<sample>.<TYPE>.junction_proof.tsv     # seam recovery proof
  consensus/consensus_qc.tsv                       # aggregated QC
  coverage/                                      # per-amplicon / per-base depth
  figures/                                         # SVG + PDF + 300-dpi PNG
  run_summary.tsv                                  # sample -> types present
```

Chapter 9

Interpreting outputs

9.1 run_summary.tsv

One row per sample: `sample_id` and the `;`-separated `types_present`. From the validation run:

```
sample_id       types_present
SAMPLE_HPV16    HPV16
SAMPLE_HPV58    HPV58
SAMPLE_COINF    HPV16;HPV58
```

9.2 consensus_qc.tsv

One row per (sample × type):

column	meaning
`sample_id`, `type`	the sample and HPV type
`pct_genome`	% of the reference reconstructed (non-N)
`pct_N`	% of the consensus that is N
`engine`	caller path used (`artic+clair3` or `portable`)
`model`	Clair3 model (e.g. `r1041_e82_400bps_hac_g615`)
`seam_N`	N bases across the origin **after** the junction merge
`junction_engine`	engine used for the amp3 junction run
`junction_bases_filled`	seam bases recovered from the rotated amp3 run

9.3 Consensus FASTAs / VCFs

- **…**`consensus.fasta` — the per-type whole-genome consensus (origin-continuous after merge).
- **…**`pass.vcf` / **…**`merged.vcf.gz` — Clair3 variant calls (per ARTIC read-group, then merged).

9.4 The QC figures (`results/figures/`)

Each is emitted as SVG + PDF + 300-dpi PNG:

- **`03_amplicon_heatmap_<TYPE>`** — per-amplicon coverage heatmap (sample × amplicon, log-depth); the single most useful view (dropout shows as a cold cell).
- **`04_depth_<sample>_<TYPE>`** — per-base depth track (even tiling, seam continuity).
- **`01_read_counts_yield`** — reads/yield per sample.
- **`02_length_quality`** — length × quality (confirm the ~1.1 kb HPV52 amp survives).
- **`06_type_assignment`** — binning per sample (co-infections show as multiple bars).
- **`07_consensus_completeness`** — % genome / % N per sample×type.
- Pool-balance panel — Pool 1 vs Pool 2 representation (needs both pools).

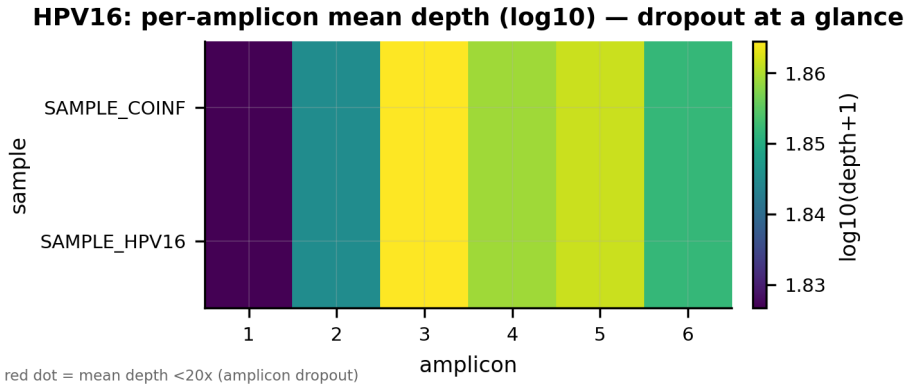


Figure 9.1: Per-amplicon coverage heatmap (HPV16) — log-depth per amplicon; cold cells = dropout.

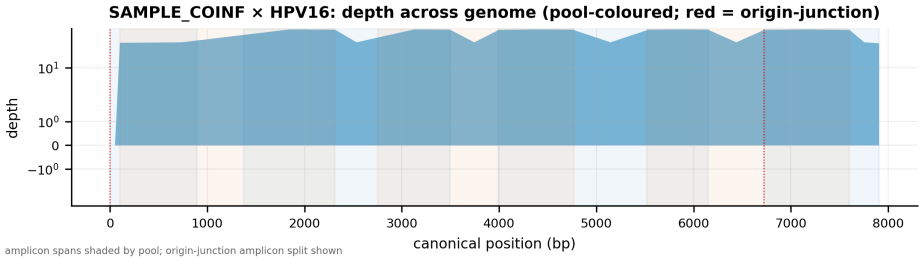


Figure 9.2: Per-base depth track, SAMPLE_COINF / HPV16 — even tiling and origin continuity.

9.5 Design caveats analysts must expect

- **HPV52 amp7 over-representation.** The short ~1.1 kb L1 gap-fill amplicon over-amplifies in multiplex, so HPV52 depth is skewed toward amp7. Expected; the wet-lab mitigation is to reduce HPV52_7 primer concentration or post-PCR normalise.
- **HPV35 amp4 lower depth on divergent isolates.** HPV35 amp4 has a documented, African-justified inclusivity exception (panel inclusivity 75 % from

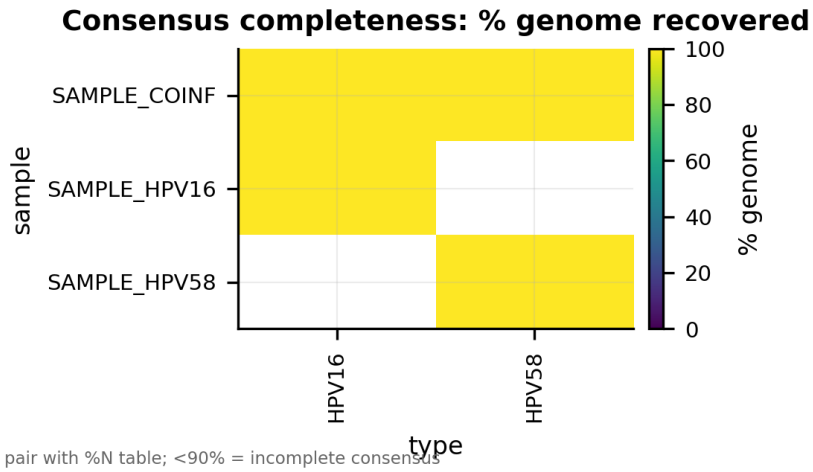


Figure 9.3: Consensus completeness — % genome reconstructed and % N per sample×type.

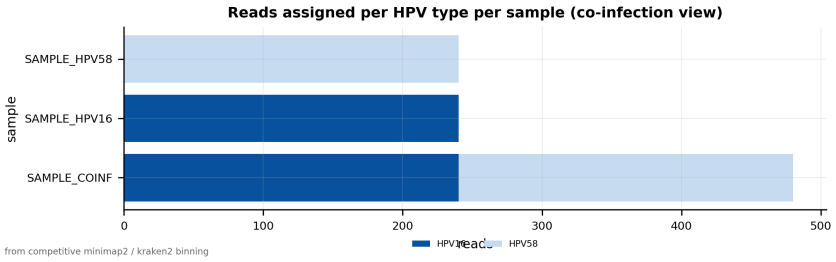


Figure 9.4: Type assignment — reads binned per HPV type; co-infection shows two types in one sample.

USA/divergent isolates; African inclusivity 100 %). Expect lower amp4 depth on non-African divergent isolates — not a pipeline fault.

Chapter 10

Validation record

Two independent software validations were performed. Both are software validations; **biological validation on control material is the next, separate, wet-lab step.**

10.1 (1) End-to-end on the real ARTIC + Clair3 path

A built-in synthetic dataset of three samples — `SAMPLE_HPV16`, `SAMPLE_HPV58`, and `SAMPLE_COINF` (a deliberate HPV16 + HPV58 co-infection) — was run end to end through the real ARTIC + Clair3 path (`--use-artic --require-artic`) on the cluster (16 threads). From the actual `consensus_qc.tsv`:

sample_id	type	pct_genome	pct_N	engine	model	seam_N	junction_engine	filled
SAMPLE_HPV16	HPV16	100.0	0.0	artic+clair3	r1041_e82_400bps_hac_g615	0	artic+clair3	448
SAMPLE_HPV58	HPV58	100.0	0.0	artic+clair3	r1041_e82_400bps_hac_g615	0	artic+clair3	588
SAMPLE_COINF	HPV16	100.0	0.0	artic+clair3	r1041_e82_400bps_hac_g615	0	artic+clair3	448
SAMPLE_COINF	HPV58	100.0	0.0	artic+clair3	r1041_e82_400bps_hac_g615	0	artic+clair3	588

Columns: `pct_genome` / `pct_N` = % genome reconstructed / % N; `seam_N` = origin N after merge; `junction_engine` = caller for the amp3 run; `filled` = `junction_bases_filled`.

Evidence this is the real caller path (not the portable fallback): the ARTIC `minion.log.txt` shows `align_trim --normalise 100 … --primer-match-threshold 35 --min-mapq 20` on the `no_amp3` BED (main run) and the rotated BED (junction run), and `run_clair3.sh … --platform=ont --model_path=…/r1041_e82_400bps_hac_g615` producing the per-amplicon VCFs.

Conclusions: all four sample×type combinations reconstruct **100 % of the genome, 0 % N**; the origin seam is fully recovered (`merged_seam_N = 0` at 100 % identity, via the real ARTIC+Clair3 two-run mechanism — see Origin-junction handling); and the co-infection is correctly resolved (`SAMPLE_COINF -> HPV16;HPV58`, both reconstructed completely).

10.2 (2) All-users deployment verified under a non-admin account

`verify_env.sh` was run by a genuine **non-admin** user on a node: **53 passed, 0 failed.** It confirmed:

- **Named-env activation** — all six `hpv-*` envs resolve and activate by name (`conda activate hpv-<name>`), with clean tool-version rows.
- **Deployed pipeline end-to-end** — `simulate_test_data.py` + `run_pipeline.py` from `$HPV_PIPELINE` ran to completion, producing `run_summary.tsv`, the co-infection split, figures, and a valid sorted+indexed BAM (`samtools quickcheck` on a real `*.canonical.bam`).
- **No writes to the shared install** — the non-admin user could not write `/hpc/opt/hpv-bioinfo`, the shared conda, or `$HPV_PIPELINE`; all outputs landed in the user's own workdir/HOME.
- **Silent profile auto-load** — `/etc/profile.d/hpv-bioinfo.sh` sources with zero output, sets `HPV_PIPELINE` and puts `hpv-run` on PATH, runs no conda init, and does not duplicate PATH on re-source.
- **Clair3 2.x models** — fast/hac/sup present, readable, PyTorch (2.x) format.

Chapter 11

Troubleshooting & FAQ

`CommandNotFoundError: … hpv-artic` / env not found. Run `conda env list` — you should see `hpv-qc … hpv-smk` next to `wgs_*`. If missing, the shared install isn't on this node's conda; contact an admin. Never create the envs yourself.

`hpv-run: command not found`. Normally `hpv-run` is on PATH automatically at login. If you are in a non-login/batch shell, `source /etc/profile.d/hpv-bioinfo.sh` first. Confirm with `echo "$HPV_PIPELINE"` (-> `/hpc/opt/hpv-bioinfo/pipeline`) and `command -v hpv-run`. You never need a home-directory copy; if `$HPV_PIPELINE` is missing entirely, the admin has not run the deploy step.

Clair3 "model not found". `ls "$MODELS_DIR"` to see the installed `r1041_e82_400bps_{fast,hac,sup}_*` builds. Pass an explicit `--model <name> --model-dir "$MODELS_DIR"`. Confirm the FASTQ `basecall_model_version_id` matches an installed tier.

`align_trim`: "negative dimensions are not allowed" (historical — fixed). This was the origin-junction amplicon (amp3) wrapping position 0 -> negative length. It is resolved by the two-run design (see Origin-junction handling). If you see it again, you are pointing ARTIC at a raw canonical BED that still contains amp3 — use the pipeline's consensus stage, which builds the `no_amp3` and rotated BEDs for you.

Low amplicon depth. Check the `03_amplicon_heatmap` figure. Expected design hotspots: HPV52 amp7 (high) and HPV35 amp4 on divergent non-African isolates (low) — see Design caveats. Otherwise suspect input depth or a pool imbalance (confirm both pools were included).

Co-infection: only one type expected but two reported. The binner is co-infection-aware by design. Check `typing/<sample>.types.tsv` read counts; raise `--min-reads` if a low-count second type is index-hopping noise rather than a true co-infection.

A sample×type failed and the run aborted. That is `--require-artic` doing its job (strict, no silent fallback). Read `results/consensus/<sample>.<TYPE>.artic.main/*.minion.log.txt`. For quick triage you may re-run with `--allow-fallback`, but **do not report fallback consensus as ARTIC+Clair3 results**.

Where do logs/outputs go? Under your `--outdir` (default `./results`); per-stage ARTIC logs in `results/consensus/<sample>.<TYPE>.artic.{main,rotj}/`. Nothing is written to `/hpc/opt`.

Who do I contact for admin changes? Anything touching `/hpc/opt/hpv-bioinfo` or the shared conda (new models, scheme updates, env rebuilds, pipeline updates) is admin-only — file a request with the bioinformatics core.

Chapter 12

Appendix

12.1 Core runtime version strings

From each tool's own `--version` banner on the cluster (complements the conda-package versions in the env tables): artic 1.10.1 · clair3 2.0.1 · minimap2 2.31-r1302 · samtools 1.23.1 · bcftools 1.23.1 · nextflow 26.04.3 · snakemake 9.22.0 · python 3.13.13.

12.2 Full `conda list` per env

The complete package list for all six envs (every dependency, version, build string, channel) is shipped beside this manual as **`hpv_env_versions.txt`**, captured live from the production cluster. Headline summary:

Env	Python	Headline installed versions
`hpv-qc`	3.13.13	nanoplot 1.46.2 · nanocomp 1.25.6 · pycoqc 2.5.0.3 · multiqc 1.35 · fastqc 0.12.1 · seqkit 2.13.0 · csvtk 0.37.0 · mosdepth 0.3.14 · bedtools 2.31.1 · samtools 1.23.1 · chopper 0.13.0
`hpv-artic`	3.11.15	**artic 1.10.1 · clair3 2.0.1** · pytorch-cpu 2.10.0 · align_trim 1.1.0 · primalbedtools 1.0.0 · minimap2 2.31 · samtools/bcftools/htslib 1.23.1 · whatshap 2.8 · longphase 2.0.2
`hpv-asm`	3.12.13	flye 2.9.6 · **medaka 2.2.2** · racon 1.5.0 · minimap2 2.31 · samtools 1.23.1 · pytorch-cpu 2.9.1 · any2fasta 0.8.1
`hpv-py`	3.13.13	pysam 0.24.0 · pandas 3.0.3 · numpy 2.4.6 · scipy 1.17.1 · scikit-bio 0.7.3 · matplotlib 3.10.9 · seaborn 0.13.2 · plotly 6.8.0 · kaleido 1.3.0 · minimap2 2.31 · mosdepth 0.3.14
`hpv-typing`	3.14.5	minimap2 2.31 · samtools 1.23.1 · kraken2 2.17.1 · seqkit 2.13.0 · blast 2.17.0 · entrez-direct 24.0
`hpv-smk`	3.13.13	**snakemake-minimal 9.22.0** · pandas 3.0.3
`bin/`	—	nextflow 26.04.3 (runtime) + the `hpv-run` launcher

Regenerate the capture on the node:

```
for e in hpv-qc hpv-artic hpv-asm hpv-py hpv-typing hpv-smk; do
  echo "=== $e ==="; conda list -n "$e"
done > hpv_env_versions.txt
```

The exact resolved builds are also exported at install time to `/hpc/opt/hpv-bioinfo/envs_locks/<env>.loc` (build strings in `.fullbuild.lock.yml`).

12.3 The all-users permissions model

- Shared install `INSTALL_ROOT=/hpc/opt/hpv-bioinfo` owned `root:bioinfo`, dirs 2775 (setgid), files world-readable — read/execute for all users on any node.
- Named `hpv-*` envs in `/hpc/opt/conda/envs` (created like the lab's `wgs_*` envs).
- No analysis step writes to the shared install. Run-time output/cache/work/temp go to the user's CWD/HOME (`NXF_HOME`, `APPTAINER_CACHEDIR`, `CONDA_PKGS_DIRS` -> `$HOME`). Clair3 models are read-only, passed by `--model-dir`.
- Only admins modify the shared install; `setup_hpv_nanopore_env.sh` is idempotent and name-scoped (only ever creates/rebuilds `hpv-*`, never `base`/`wgs_*`).

12.4 Design provenance

- **Scheme version v5**, 12 high-risk types, 6 amplicons each, dual-pool (Pool 1 / Pool 2), one origin-spanning junction amplicon (amp3) per type. Per-type clade, canonical accession, genome length, junction amplicon and degenerate-primer counts are in `scheme_summary.tsv` (e.g. HPV16/Alpha9/K02718/7906 bp; HPV58/Alpha9/D90400/7824 bp).
- Documented design exceptions: HPV35 amp4 African-justified inclusivity exception; HPV52 amp7 ~1.1 kb L1 gap-fill over-amplification (`softsite_exceptions.md`, `wet-lab_balancing_notes.md`).
- Full design methods & phylogeny: see the project METHODS_REPORT / `FINAL_REPORT_v1-v5`.

12.5 How to report issues

Include: the exact command, the relevant `results/consensus/<sample>.<TYPE>.artic.*/…minion.log.t` the `consensus_qc.tsv` / `junction_proof.tsv` rows, and `conda list -n hpv-artic`. Send to the bioinformatics core. For shared-install changes (models, schemes, envs, pipeline updates), file an admin request — do not modify `/hpc/opt` yourself.

End of document. Manual v1.1 · 2026-06-06 · generated from the live `conda list` (`hpv_env_versions.txt`), the real ARTIC+Clair3 synthetic validation run, and the non-admin `verify_env.sh` result (53 passed, 0 failed).